

# COMP 4736

# Introduction to Operating Systems

## 01. Computer Architecture and Overview

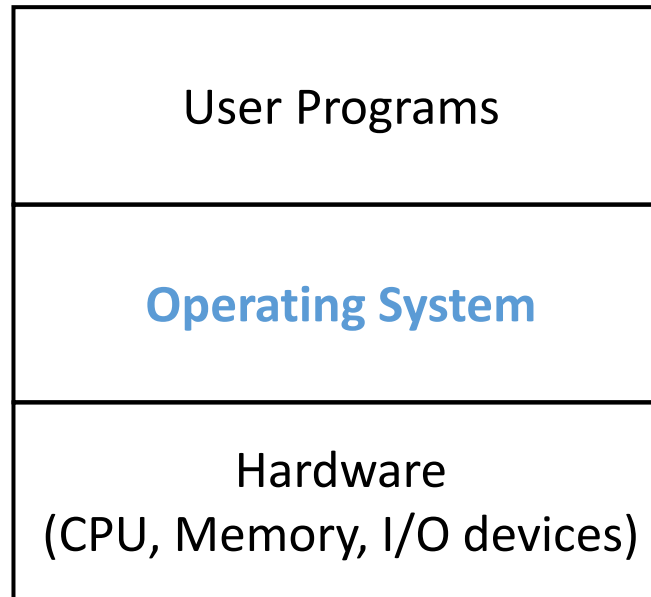
(OSTEP Ch. 2 & 4)

# Ch. 02. Introduction to Operating Systems

(OSTEP Ch 2)

# What is an Operating System?

- Middleware between **user programs** and **system hardware**.
- Manages hardware: CPU, main memory, I/O devices (disk, network card, mouse, keyboard, etc.)



# What happens when a program runs?

(Recall: a program is a sequence of instructions and data.)

- The processor **fetches** an instruction from memory, ...
- **decodes** it (i.e., figures out which instruction this is), ...
- and **executes** it (e.g., add two numbers, access memory, check a condition, jump to a function, etc.).
- The processor then moves to the **next instructions**. It repeats until program completes.

# Operating Systems (OS)

- Responsible for
  - Making it easy to **run** programs.
  - Allowing programs to **share** memory.
  - Enabling programs to **interact** with devices.

OS is in charge of making sure the system operates **correctly** and **efficiently**.

# Virtualization

- The OS takes a **physical** resource and transforms it into a **virtual** form of itself.
  - Physical resource: processor, memory, disk, etc.
- The virtual form is more **general**, **powerful** and **easy-to-use**.
- Sometimes, we refer to the OS as a **virtual machine**.

# System Call

- **System call** allows user to **tell the OS** what to do.
- The OS provides some interfaces (**APIs, standard library**).
- A typical OS exports a few hundred **system calls**.
  - Run programs
  - Access memory
  - Access devices

# OS as a Resource Manager

- The OS **manages resources** such as CPU, memory and disk.
- The OS allows
  - Many programs to run → Sharing the **CPU**
  - Many programs to concurrently access their own instructions and data → Sharing **memory**
  - Many programs to access devices → Sharing **disks**



# Virtualizing the CPU

- The system has a very large number of virtual CPUs.
- Turning a single CPU into a **seemingly infinite number** of CPUs.
- Allowing many programs to **seemingly run at once**  
→ **Virtualizing the CPU**

# Virtualizing the CPU (cpu.c)

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <sys/time.h>
4  #include <assert.h>
5  #include "common.h"
6
7  int
8  main(int argc, char *argv[])
9  {
10     if (argc != 2) {
11         fprintf(stderr, "usage: cpu <string>\n");
12         exit(1);
13     }
14     char *str = argv[1];
15     while (1) {
16         Spin(1);
17         printf("%s\n", str);
18     }
19     return 0;
20 }
```

# Virtualizing the CPU (Result 1)

```
prompt> gcc -o cpu cpu.c -Wall
prompt> ./cpu "A"
A
A
A
A
^C
prompt>
```

- Runs forever.
- We halt the program by pressing Ctrl+C.

# Virtualizing the CPU (Result 2)

```
prompt> ./cpu A & ./cpu B & ./cpu C & ./cpu D &  
[1] 7353  
[2] 7354  
[3] 7355  
[4] 7356  
A  
B  
D  
C  
A  
B  
D  
C  
A  
...
```

Even though we have only **one processor**, all **four** of **programs** seem to be **running at the same time!**

# Virtualizing the CPU

- Running multiple programs at once raises new questions.
  - E.g., if two programs want to run at the same time, which **should** run?
- It is handled by **policies** of the OS.
- We will study them as we learn about the basic **mechanisms** that OS implement, as a **resource manager**.

# Virtualizing Memory

- The **physical memory** is an array of bytes.
- A program keeps all of its data structures in memory.
- **Read** memory (load):
  - Specify an address to be able to access the data
- **Write** memory (store):
  - Specify the data to be written to the given address

# Virtualizing Memory (mem.c)

```
1  #include <unistd.h>
2  #include <stdio.h>
3  #include <stdlib.h>
4  #include "common.h"
5
6  int
7  main(int argc, char *argv[])
8  {
9      int *p = malloc(sizeof(int));           // a1
10     assert(p != NULL);
11     printf("(%d) address pointed to by p: %p\n",
12            getpid(), p);                     // a2
13     *p = 0;                                  // a3
14     while (1) {
15         Spin(1);
16         *p = *p + 1;
17         printf("(%d) p: %d\n", getpid(), *p); // a4
18     }
19     return 0;
20 }
```

# Virtualizing Memory (Result 1)

```
prompt> ./mem
(2134) address pointed to by p: 0x200000
(2134) p: 1
(2134) p: 2
(2134) p: 3
(2134) p: 4
(2134) p: 5
^C
```

- The newly allocated memory is at address 0x200000.
- It updates the value and prints out the result.



# Virtualizing Memory (Result 2)

```
prompt> ./mem & ./mem &  
[1] 24113  
[2] 24114  
(24113) address pointed to by p: 0x200000  
(24114) address pointed to by p: 0x200000  
(24113) p: 1  
(24114) p: 1  
(24114) p: 2  
(24113) p: 2  
(24113) p: 3  
(24114) p: 3  
(24113) p: 4  
(24114) p: 4
```

- It is as if each running program has its **own private memory**.
  - Each running program has allocated memory **at the same address**.
  - Each seems to be updating the value at 0x200000 independently.

# Virtualizing Memory

- Each process accesses its own private **virtual address space**.
- The OS maps **address space** onto the **physical memory**.
- A memory reference within one running program **does not affect** the address space of other processes.
- Physical memory is a **shared resource**, managed by the OS.

# Problems of Concurrency

- The OS is juggling **many things at once**, first running one process, then another, and so forth.
- Modern **multi-threaded programs** also exhibit the concurrency problem.

# Concurrency (thread.c)

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include "common.h"
4  #include "common_threads.h"
5
6  volatile int counter = 0;
7  int loops;
8
9  void *worker(void *arg) {
10     int i;
11     for (i = 0; i < loops; i++) {
12         counter++;
13     }
14     return NULL;
15 }
16
17 int main(int argc, char *argv[]) {
18     if (argc != 2) {
19         fprintf(stderr, "usage: threads <value>\n");
20         exit(1);
21     }
```

# Concurrency (thread.c)

```
22     loops = atoi(argv[1]);
23     pthread_t p1, p2;
24     printf("Initial value : %d\n", counter);
25
26     Pthread_create(&p1, NULL, worker, NULL);
27     Pthread_create(&p2, NULL, worker, NULL);
28     Pthread_join(p1, NULL);
29     Pthread_join(p2, NULL);
30     printf("Final value : %d\n", counter);
31     return 0;
32 }
```

- The main program creates **two threads**.
  - **Thread**: a function running within the **same memory space**. Each thread start running in a routine called `worker()`.
  - `worker()`: increments a counter

# Concurrency (Result)

- loops determines how many times each of the two workers will increment the shared counter in a loop.
  - loops: 1000.

```
prompt> gcc -o thread thread.c -Wall -pthread
prompt> ./thread 1000
Initial value : 0
Final value : 2000
```

- loops: 100000.

```
prompt> ./thread 100000
Initial value : 0
Final value : 143012           // huh??
prompt> ./thread 100000
Initial value : 0
Final value : 137298           // what the??
```

# Why is this happening?

- Increment a shared counter → take three instructions.
  1. Load the value of the counter from memory into register.
  2. Increment it.
  3. Store it back into memory.
- These three instructions do not execute **atomically**.
  - Problem of **concurrency** can happen.

# Persistence

- Devices such as DRAM store values in a **volatile**.
- **Hardware** and **software** are needed to store data **persistently**.
- Hardware: I/O device such as a **hard drive**, **solid-state drives (SSDs)**
- Software:
  - **File system** manages the disk.
  - File system is responsible for storing any **files** the user creates.



# Persistence (io.c)

```
1  #include <stdio.h>
2  #include <unistd.h>
3  #include <assert.h>
4  #include <fcntl.h>
5  #include <sys/types.h>
6
7  int main(int argc, char *argv[]) {
8      int fd = open("/tmp/file", O_WRONLY|O_CREAT|O_TRUNC,
9          S_IRWXU);
10     assert(fd > -1);
11     int rc = write(fd, "hello world\n", 13);
12     assert(rc == 13);
13     close(fd);
14     return 0;
15 }
```

- `open()`, `write()`, and `close()` **system calls** are routed to the part of OS called the **file system**, which handles the requests.

# Persistence

- What OS does in order to write to disk?
  - Figure out **where** on disk this new data will reside.
  - **Issue I/O** requests to the underlying storage device.
- File system handles system crashes during write.
  - **Journaling** or **copy-on-write**.
  - Carefully **ordering** writes to disk.

# Design Goals

- Build up **abstraction**
  - Make the system convenient and easy to use.
- Provide high **performance**
  - **Minimize the overheads** of the OS.
  - OS must strive to provide virtualization without excessive overheads.
- **Protection** between applications
  - **Isolation**: Bad behavior of one does not harm other and the OS itself.

# Design Goals

- High degree of **reliability**
  - The OS must also run non-stop.
- Other issues
  - **Energy-efficiency**
  - **Security**
  - **Mobility**

# Some History

- IBM 701 — Electronic Data Processing Machine (1952)
- No OS. Operated manually.



# Batch Processing

- IBM 709 —Data Processing System (1952)
- SHARE Operation System (SOS). Manages buffers and I/O devices.



# Multiprogramming

- PDP-11 — minicomputer (1970)
- Unix. Supports multiprogramming.



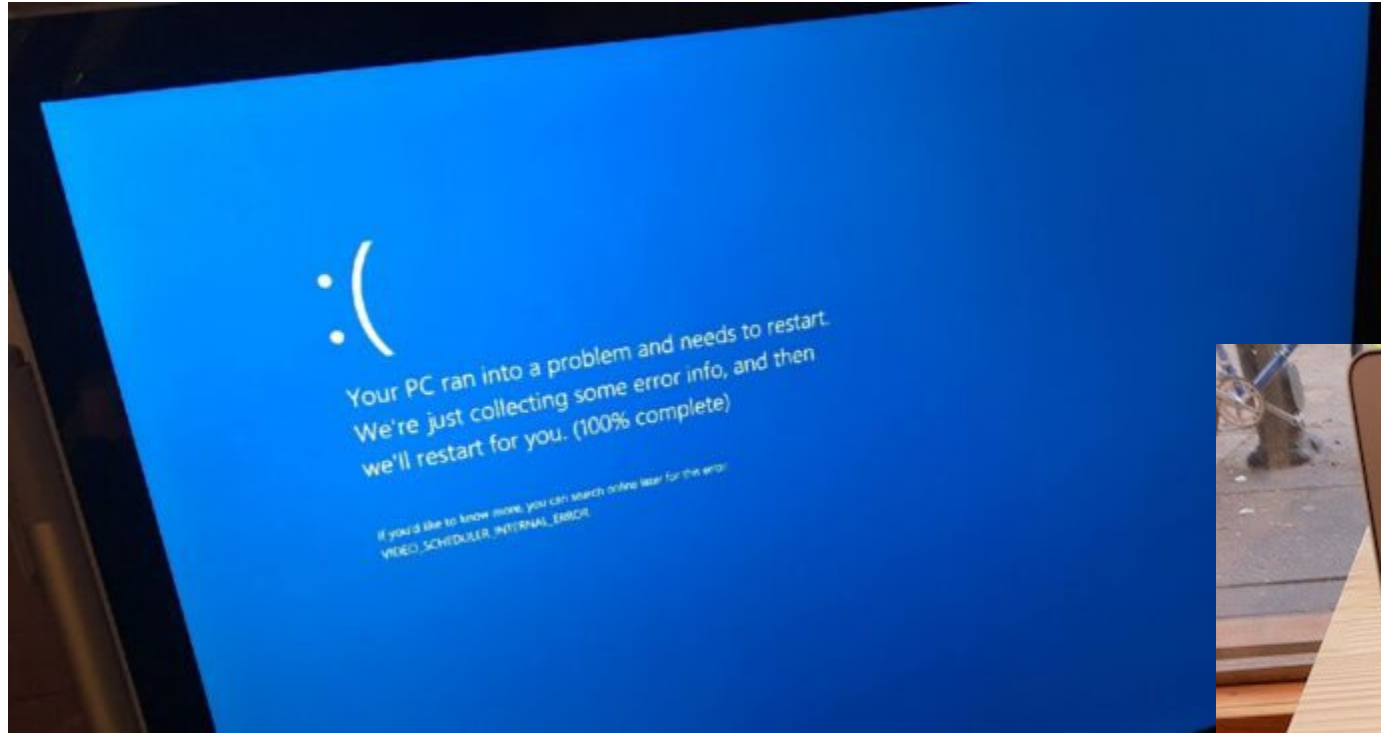
# Personal Computer (PC)

- Apple II (1977) and IBM PC (1981)
- Disk Operating System (DOS) and Mac OS





# Modern Era (and Error)



# Mobile and IoT



# Summary

- We have briefly introduced what an **operating system (OS)** is.
- Included in this course:
  - **Virtualization** of CPU and memory
    - Resource management. Allows multiple programs to run.
  - **Concurrency**
    - Data consistency. Data protection.
- Not included (but important topics nonetheless):
  - **Persistence** via devices and file systems.
  - **Networking**
  - **Graphics devices**
  - **Security**

# References

- P. Brinch Hansen, “The Evolution of Operating Systems”, *Classic Operating Systems: From Batch Processing to Distributed Systems*, Springer-Verlag, New York, 2000.
- IBM Archives: IBM 701.  
[https://www.ibm.com/ibm/history/exhibits/701/701\\_intro.html](https://www.ibm.com/ibm/history/exhibits/701/701_intro.html)
- IBM Archives: 709 Data Processing System  
[https://www.ibm.com/ibm/history/exhibits/mainframe/mainframe\\_PP709.html](https://www.ibm.com/ibm/history/exhibits/mainframe/mainframe_PP709.html)
- A, Hudson, “A brief tour of the PDP-11, the most influential minicomputer of all time”, Ars Technica, Mar 2022.  
<https://arstechnica.com/gadgets/2022/03/a-brief-tour-of-the-pdp-11-the-most-influential-minicomputer-of-all-time/7/>

# Ch. 04. The Process

(OSTEP Ch. 4)

# How to Provide the Illusion of Many CPUs?

- **CPU virtualization**
  - The OS can promote the illusion that many virtual CPUs exist.
  - **Time sharing**: Running one process, then stopping it and running another.
- The potential cost is **performance**.

# Implementing Virtualization

- OS will need low-level machinery and high-level intelligence.
- **Mechanisms**: How to “context switch” between processes.
- **Policies**: Decides which process to run, e.g., a **scheduling policy**.

# Process

A process is a **running program**.

- Comprising of a process:
- **Memory (address space)**
  - Instructions
  - Data section
- **Registers**
  - Program counter
  - Stack pointer



# Process API

- These APIs are available on any modern OS.
- **Create**
  - Create a new process to run a program
- **Destroy**
  - Halt a runaway process
- **Wait**
  - Wait for a process to stop running
- **Miscellaneous Control**
  - Some kind of method to suspend a process and then resume it
- **Status**
  - Get some status info about a process

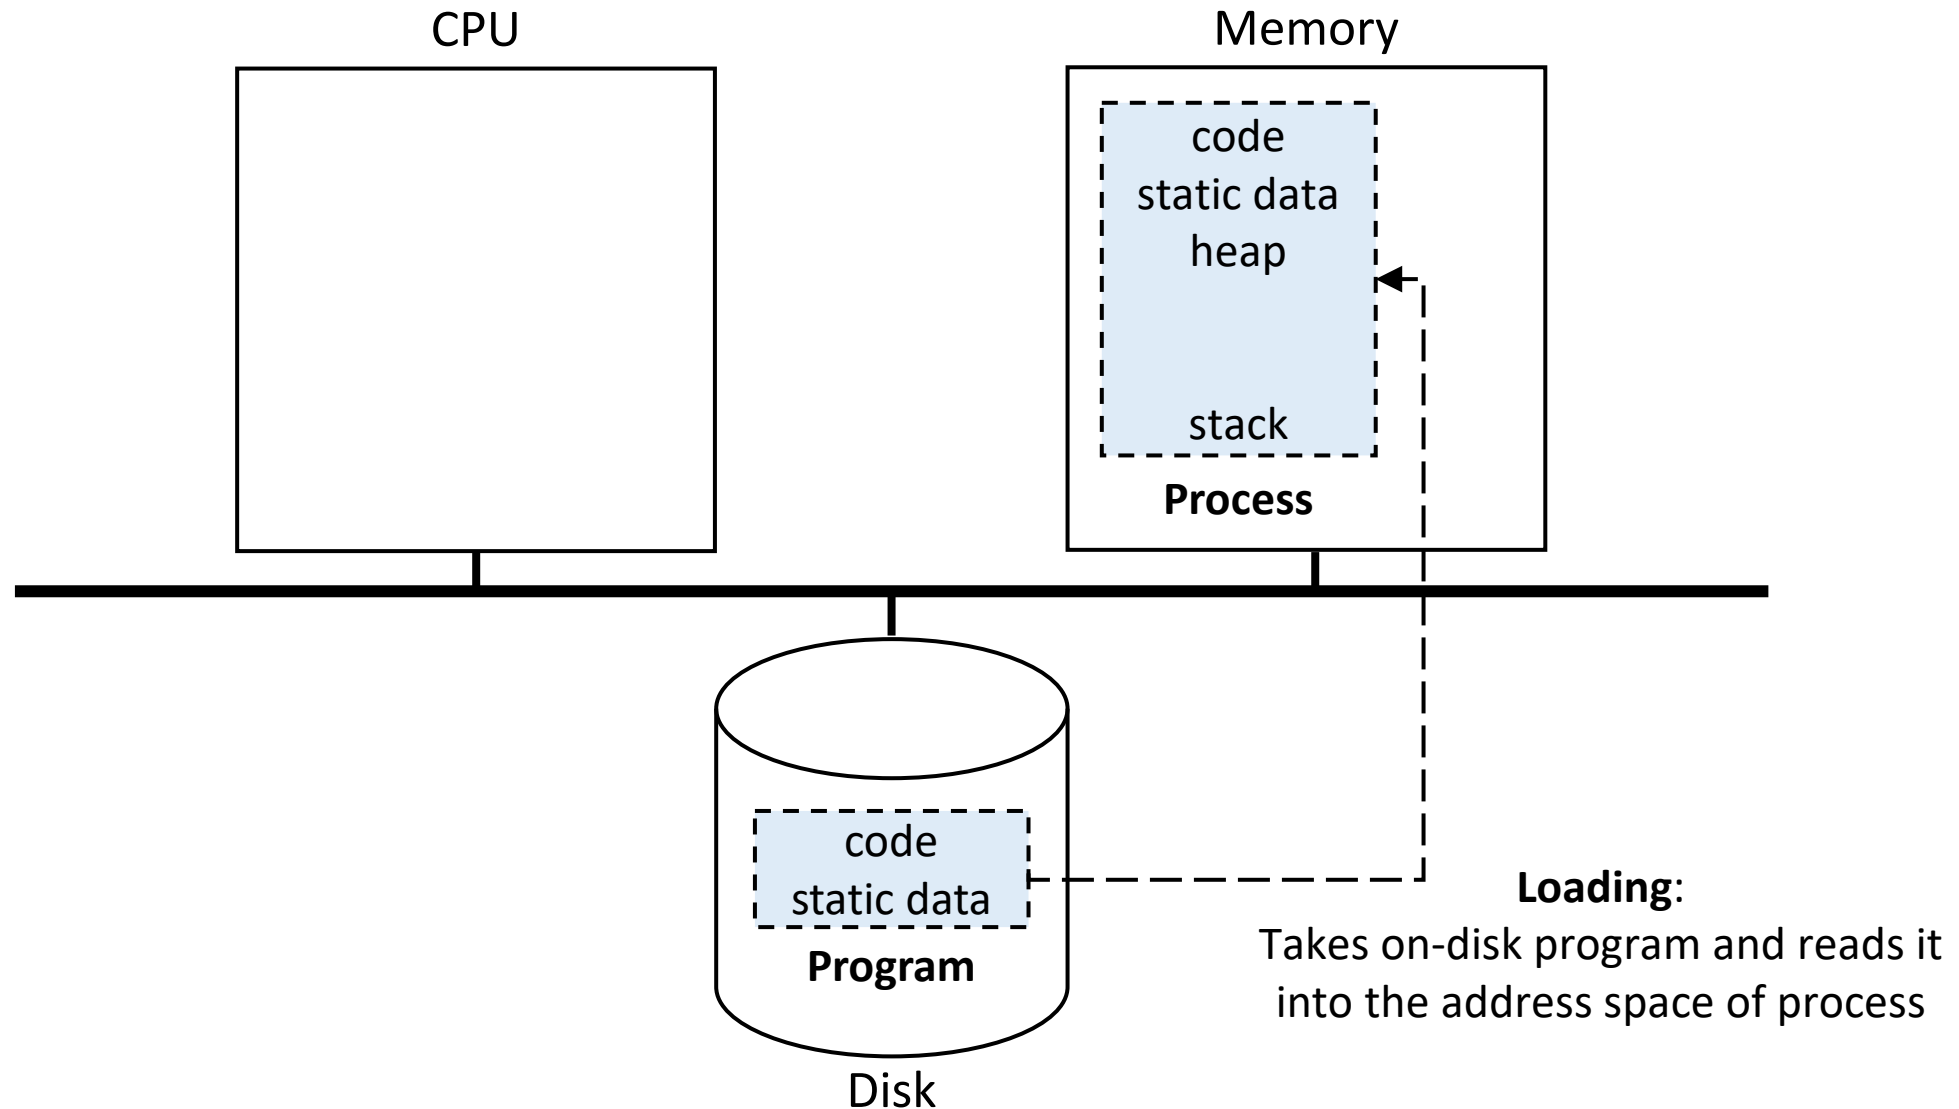
# Process Creation

1. **Load** a program code into **memory**, into the address space of the process.
  - Programs initially reside on **disk** in **executable format**.
  - Modern OS perform the loading process **lazily**.
    - Loading pieces of code or data **only** as they are needed during program execution.
2. The program's run-time **stack** is allocated.
  - Use the stack for local variables, function parameters, and return address.
  - Initialize the stack with arguments → argc and the argv array of `main()` function

# Process Creation

3. The program's **heap** is created.
  - Used for explicitly requested dynamically allocated data.
  - Program request such space by calling `malloc()` and free it by calling `free()`.
4. The OS do some other **initialization** tasks.
  - Input/output (I/O) setup
  - Each process by default has three open file descriptors.
    - Standard input, output and error
5. **Start the program** running at the entry point, namely `main()`.
  - The OS **transfers control** of the CPU to the newly-created process.

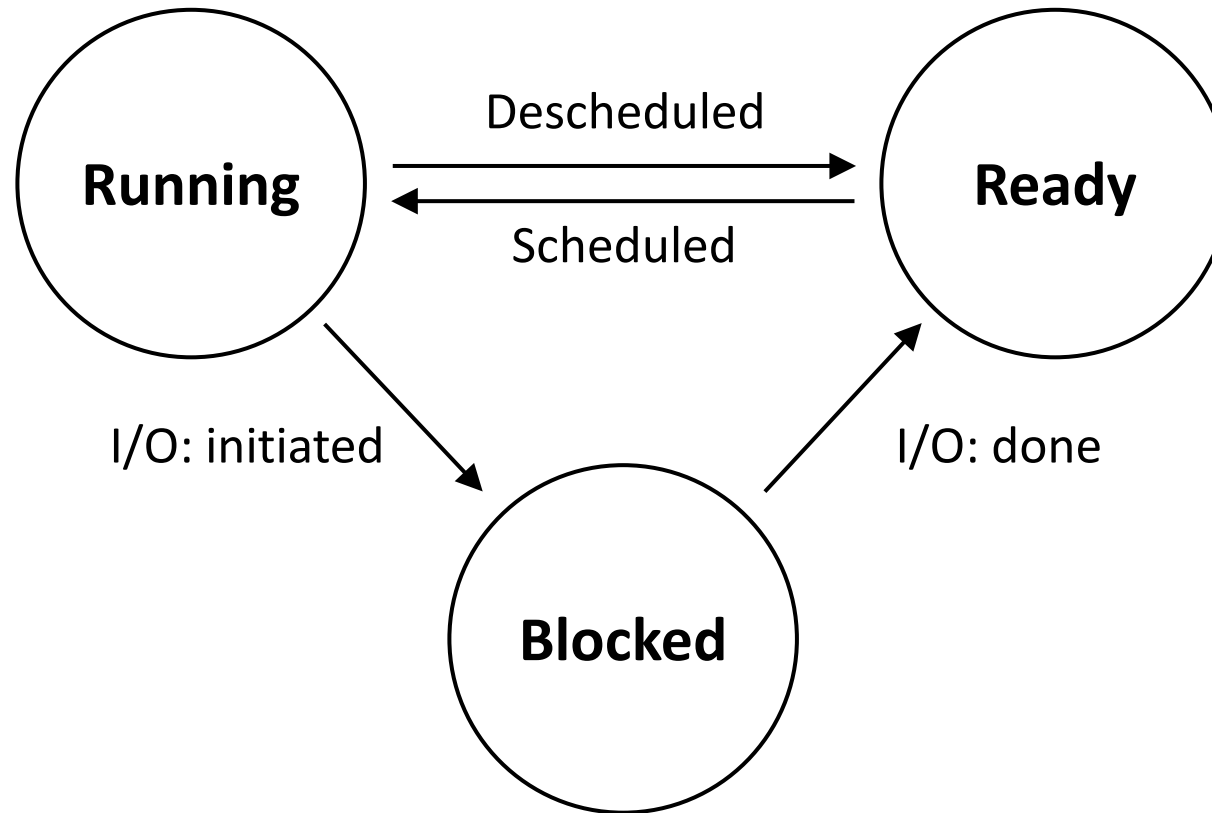
# Loading: From Program To Process



# Process States

- A process can be one of three states.
- **Running**
  - A process is running on a processor.
- **Ready**
  - A process is ready to run but for some reason the OS has chosen not to run it at this given moment.
- **Blocked**
  - A process has performed some kind of operation.
  - When a process initiates an I/O request to a disk, it becomes blocked and thus some other process can use the processor.

# Process: State Transitions



# Tracing Process State: CPU Only

Time	Process <sub>0</sub>	Process <sub>1</sub>	Notes
1	Running	Ready	
2	Running	Ready	
3	Running	Ready	
4	Running	Ready	Process <sub>0</sub> is now done
5	—	Running	
6	—	Running	
7	—	Running	
8	—	Running	Process <sub>1</sub> is now done

# Tracing Process State: CPU and I/O

Time	Process <sub>0</sub>	Process <sub>1</sub>	Notes
1	Running	Ready	
2	Running	Ready	
3	Running	Ready	Process <sub>0</sub> initiates I/O
4	Blocked	Running	Process <sub>0</sub> is blocked, so Process <sub>1</sub> runs
5	Blocked	Running	
6	Blocked	Running	
7	Ready	Running	I/O done
8	Ready	Running	Process <sub>1</sub> is now done
9	Running	—	
10	Running	—	Process <sub>0</sub> is now done



# Data Structures

- OS maintains a data structure (e.g., **process list**) of all ready, running and blocked processes.
- **Process Control Block (PCB)**
  - A C-structure that contains information about each process.
  - **Register context**: a set of registers that define the state of a process

# The xv6 Proc Structure

```
// the information xv6 tracks about each process
// including its register context and state
struct proc {
    char *mem;           // Start of process memory
    uint sz;             // Size of process memory
    char *kstack;        // Bottom of kernel stack for this process
    enum proc_state state; // Process state
    int pid;             // Process ID
    struct proc *parent;  // Parent process
    void *chan;          // If !zero, sleeping on chan
    int killed;          // If !zero, has been killed
    struct file *ofile[NOFILE]; // Open files
    struct inode *cwd;    // Current directory
    struct context context; // Switch here to run process
    struct trapframe *tf; // Trap frame for the current interrupt
};
```

# The xv6 Proc Structure

```
// the registers xv6 will save and restore
// to stop and subsequently restart a process
struct context {
    int eip;
    int esp;
    int ebx;
    int ecx;
    int edx;
    int esi;
    int edi;
    int ebp;
};

// the different states a process can be in
enum proc_state { UNUSED, EMBRYO, SLEEPING, RUNNABLE, RUNNING, ZOMBIE };
```

# Summary

- We have introduced the most basic abstraction of the OS: the **process**.
- Low-level **mechanisms** are needed to switch between processes.
- High-level **policies** are required to schedule the processes in an intelligent way.